
CS 301

High-Performance Computing

Lab 1 - Q3A

Problem C-1: QUICK_SORT

Aditya Nawal (202001402)
Divya Patel (202001420)

February 8, 2023

Contents

1	Introduction	3
2	Hardware Details	3
2.1	Lab 207 PC	3
2.2	HPC Cluster	4
3	Problem C1 - Quicksort Algorithm	5
3.1	Description of the problem	5
3.2	Serial Complexity	5
3.3	Profiling Information	5
3.4	Optimization Strategy	6
3.5	Graph of Problem Size vs Runtime	6
3.5.1	Graph of Problem Size vs Runtime for LAB207 PCs	6
3.5.2	Graph of Problem Size vs Runtime for HPC Cluster	7
3.6	Discussion	7

1 Introduction

This report evaluates the performance of three different algorithms: the Trapezoidal rule, vector addition, and quick sort. The performance of each algorithm is measured in terms of its compute-to-memory access ratio, as well as its run-time and memory utilisation. The hardware specifications, compiler, and optimisation flags used are also described, along with the input parameters, output, and accuracy checks. The performance of each algorithm is then plotted against the problem size and compared to the theoretical performance of the machine where the code is executed. The report concludes with observations and comments about the results, including the effectiveness of each algorithm for each specific problem.

2 Hardware Details

2.1 Lab 207 PC

- Architecture: x86_64
- CPU op-mode(s): 32-bit, 64-bit
- Byte Order: Little Endian
- CPU(s): 4
- On-line CPU(s) list: 0-3
- Thread(s) per core: 1
- Core(s) per socket: 4
- Socket(s): 1
- NUMA node(s): 1
- Vendor ID: GenuineIntel
- CPU family: 6
- Model: 60
- Model name: Intel(R) Core(TM) i5-4590 CPU @ 3.30GHz
- Stepping: 3
- CPU MHz: 3300.000
- CPU max MHz: 3700.0000
- CPU min MHz: 800.0000
- BogoMIPS: 6585.38
- Virtualization: VT-x

- L1d cache: 32K
- L1i cache: 32K
- L2 cache: 256K
- L3 cache: 6144K
- NUMA node0 CPU(s): 0-3
- Flags: fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc arch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc aperfmperf eagerfpu pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 fma cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm epb invpcid_single tpr_shadow vnmi flexpriority ept vpid fsgsbase tsc_adjust bmi1 avx2 smep bmi2 erms invpcid xsaveopt dtherm ida arat pln pts

2.2 HPC Cluster

- Architecture: x86_64
- CPU op-mode(s): 32-bit, 64-bit
- Byte Order: Little Endian
- CPU(s): 16
- On-line CPU(s) list: 0-15
- Thread(s) per core: 1
- Core(s) per socket: 8
- Socket(s): 2
- NUMA node(s): 2
- Vendor ID: GenuineIntel
- CPU family: 6
- Model: 63
- Model name: Intel(R) Xeon(R) CPU E5-2640 v3 @ 2.60GHz
- Stepping: 2
- CPU MHz: 1976.914
- BogoMIPS: 5205.04
- Virtualization: VT-x

- L1d cache: 32K
- L1i cache: 32K
- L2 cache: 256K
- L3 cache: 20480K
- NUMA node0 CPU(s): 0-7
- NUMA node1 CPU(s): 8-15

3 Problem C1 - Quicksort Algorithm

3.1 Description of the problem

The goal of this report is to evaluate the performance of the quicksort algorithm, with a focus on its run-time. In this report, we describe the hardware specifications, compiler, and optimization flags used in the experiments, as well as the input parameters, output, and accuracy checks. Additionally, we present the results of our performance evaluation, including plots that compare the run-time of the algorithm against the problem size.

3.2 Serial Complexity

The serial complexity of the Quick Sort algorithm is $O(n \log n)$, where "n" is the size of the array.

3.3 Profiling Information

Following are the snapshots taken while profiling.

```
Flat profile:

Each sample counts as 0.01 seconds.
%   cumulative   self           self       total
time  seconds    seconds   calls   ms/call  ms/call  name
100.70    0.06    0.06         1    60.42    60.42  quick_sort.4403
  0.00    0.06    0.00         2     0.00     0.00  diff

%           the percentage of the total running time of the
time        program used by this function.

cumulative a running sum of the number of seconds accounted
seconds    for by this function and those listed above it.

self       the number of seconds accounted for by this
seconds    function alone.  This is the major sort for this
           listing.

calls      the number of times this function was invoked, if
           this function is profiled, else blank.

self       the average number of milliseconds spent in this
```

Figure 1: Profiling on Lab 207 PC using gprof

```
[202001402@gics4 lab2]$ gprof --line serial3.out gmon.out
Flat profile:

Each sample counts as 0.01 seconds.
%   cumulative   self           total
time  seconds    seconds   calls   ms/call  ms/call  name
52.94    0.10      0.10        1    100.58    191.10  integrate.4099
47.64    0.19      0.09 10000001    0.00      0.00  function.4094
0.00    0.19      0.00        2     0.00     0.00  diff

%           the percentage of the total running time of the
time        program used by this function.

cumulative  a running sum of the number of seconds accounted
seconds     for by this function and those listed above it.

self        the number of seconds accounted for by this
seconds     function alone.  This is the major sort for this
            listing.

calls       the number of times this function was invoked, if
            this function is profiled, else blank.

self        the average number of milliseconds spent in this
ms/call     function per call, if this function is profiled,
            else blank.

total       the average number of milliseconds spent in this
ms/call     function and its descendents per call, if this
            function is profiled, else blank.
```

Figure 2: Profiling on HPC Cluster using gprof

3.4 Optimization Strategy

We used median of the array as the pivot element, which results in a balanced partition and minimizes the number of comparisons. We also used loop unrolling to perform multiple additions in a single iteration which increases performance.

3.5 Graph of Problem Size vs Runtime

3.5.1 Graph of Problem Size vs Runtime for LAB207 PCs




Figure 3: Total mean execution time vs Problem size plot for optimised and non-optimised code (Hardware: LAB207 PC). Normal code takes much longer time to run than the optimised one.

3.5.2 Graph of Problem Size vs Runtime for HPC Cluster

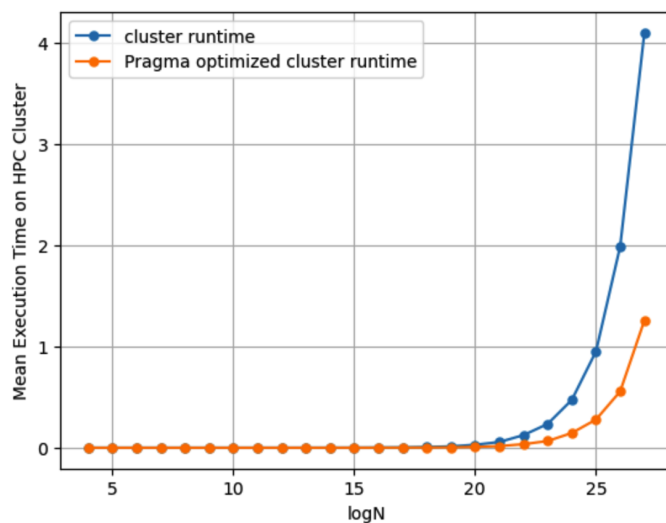


Figure 4: Total mean execution time vs Problem size plot for optimised and non-optimised code (Hardware: HPC Cluster). Normal code takes much longer time to run than the optimised one.

3.6 Discussion

The graph shows that the run-time of the algorithm increases linearly with the problem size (exponentially with respect to log of problem size). This is expected because the complexity of the algorithm is $\mathcal{O}(n \cdot \log(n))$, where "n" is the size of the arrays. The run-time of the algorithm on the HPC cluster is almost equal to the run-time on the lab PC, which is expected because though HPC cluster has more cores, our algorithm is not parallelized and thus the performance does not increase with the number of cores and single core performance is almost the same on both machines.