
CS 301

High-Performance Computing

Lab 3 - A1

Conventional Matrix Multiplication

Aditya Nawal (202001402)
Divya Patel (202001420)

February 27, 2023

Contents

1	Introduction	3
2	Hardware Details	3
2.1	Lab 207 PC	3
2.2	HPC Cluster	4
3	Problem A1	6
3.1	Description of the problem	6
3.2	Serial Complexity	6
3.3	Profiling Information	6
3.4	Optimization Strategy	7
3.5	Graph of Problem Size vs Algorithm Runtime	8
3.5.1	Graph of Problem Size vs Algorithm Runtime for LAB207 PCs	8
3.5.2	Graph of Problem Size vs Runtime for HPC Cluster	8
3.6	Graph of Problem Size vs End to End Runtime	9
3.6.1	Graph of Problem Size vs End to End Runtime for LAB207 PCs	9
3.6.2	Graph of Problem Size vs End to End Runtime for HPC Cluster	9
3.7	Discussion	10

1 Introduction

This report investigates three different approaches to matrix multiplication. In Problem A-1, we explore the conventional matrix multiplication algorithm and do it in six different ways to interchange the loops to optimize its performance. In Problem B-1, we study the use of transpose matrix multiplication as a way to reduce the computational cost of matrix multiplication. Finally, in Problem C-1, we analyze the block matrix multiplication algorithm, which uses a divide and conquer strategy to compute the product of two large matrices. For each problem, we provide a detailed analysis of the algorithms, including their computational complexity. We also implement each algorithm and evaluate its performance on matrices of different sizes. The results of our experiments are presented in the form of graphs.

2 Hardware Details

2.1 Lab 207 PC

- Architecture: x86_64
- CPU op-mode(s): 32-bit, 64-bit
- Byte Order: Little Endian
- CPU(s): 4
- On-line CPU(s) list: 0-3
- Thread(s) per core: 1
- Core(s) per socket: 4
- Socket(s): 1
- NUMA node(s): 1
- Vendor ID: GenuineIntel
- CPU family: 6
- Model: 60
- Model name: Intel(R) Core(TM) i5-4590 CPU @ 3.30GHz
- Stepping: 3
- CPU MHz: 3300.000
- CPU max MHz: 3700.0000
- CPU min MHz: 800.0000
- BogoMIPS: 6585.38

- Virtualization: VT-x
- L1d cache: 32K
- L1i cache: 32K
- L2 cache: 256K
- L3 cache: 6144K
- NUMA node0 CPU(s): 0-3
- Flags: fpv vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc arch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc aperfmperf eagerfpu pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 fma cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm epb invpcid_single tpr_shadow vnmi flexpriority ept vpid fsgsbase tsc_adjust bmi1 avx2 smep bmi2 erms invpcid xsaveopt dtherm ida arat pln pts

```
[student@localhost ~]$ ifconfig
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1 (Local Loopback)
    RX packets 60 bytes 5868 (5.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 60 bytes 5868 (5.7 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

p4p1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.100.64.86 netmask 255.255.255.0 broadcast 10.100.64.255
    inet6 fe80::b283:feff:fe97:d2f9 prefixlen 64 scopeid 0x20<link>
    ether b0:83:fe:97:d2:f9 txqueuelen 1000 (Ethernet)
    RX packets 32826 bytes 46075919 (43.9 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 8015 bytes 586362 (572.6 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

virbr0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 192.168.122.1 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:3a:16:71 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

Figure 1: IP address of Lab PC

2.2 HPC Cluster

- Architecture: x86_64
- CPU op-mode(s): 32-bit, 64-bit
- Byte Order: Little Endian
- CPU(s): 16
- On-line CPU(s) list: 0-15

- Thread(s) per core: 1
- Core(s) per socket: 8
- Socket(s): 2
- NUMA node(s): 2
- Vendor ID: GenuineIntel
- CPU family: 6
- Model: 63
- Model name: Intel(R) Xeon(R) CPU E5-2640 v3 @ 2.60GHz
- Stepping: 2
- CPU MHz: 1976.914
- BogoMIPS: 5205.04
- Virtualization: VT-x
- L1d cache: 32K
- L1i cache: 32K
- L2 cache: 256K
- L3 cache: 20480K
- NUMA node0 CPU(s): 0-7
- NUMA node1 CPU(s): 8-15

3 Problem A1

3.1 Description of the problem

Conventional matrix multiplication is a mathematical operation that takes two matrices **A** and **B** as inputs and produces a third matrix **C** as output. The output matrix is obtained by multiplying each element of a row of matrix **A** by each element of a column of matrix **B** and adding them together. The output matrix **C** has the same number of rows as matrix **A** and the same number of columns as matrix **B**. Conventional matrix multiplication is also known as standard or naive matrix multiplication because it follows the definition of matrix multiplication without any optimization or algorithmic improvement.

3.2 Serial Complexity

Conventional matrix multiplication has a time complexity of $\mathcal{O}(n^3)$ where n is the size of the matrices, which means that it becomes very slow and inefficient for large matrices.

3.3 Profiling Information

Following are the snapshots taken while profiling.

```
Flat profile:
Each sample counts as 0.01 seconds.
%   cumulative   self           self         total
time  seconds    seconds   calls   Ts/call   Ts/call   name
99.90      8.20      8.20         2      0.00     0.00    main
 0.00      8.20      0.00         2      0.00     0.00    diff

%           the percentage of the total running time of the
time        program used by this function.

cumulative  a running sum of the number of seconds accounted
seconds     for by this function and those listed above it.

self        the number of seconds accounted for by this
seconds     function alone. This is the major sort for this
            listing.

calls       the number of times this function was invoked, if
            this function is profiled, else blank.

self        the average number of milliseconds spent in this
ms/call     function per call, if this function is profiled,
            else blank.

total       the average number of milliseconds spent in this
ms/call     function and its descendents per call, if this
            function is profiled, else blank.

name        the name of the function. This is the minor sort
            for this listing. The index shows the location of
            the function in the gprof listing. If the index is
            in parenthesis it shows where it would appear in
            the gprof listing if it were to be printed.

Copyright (C) 2012 Free Software Foundation, Inc.
```

Figure 2: Screenshot of text file generated from profiling on Lab 207 PC using gprof

```

Flat profile:

Each sample counts as 0.01 seconds.
%   cumulative   self           self       total
time  seconds    seconds   calls   Ts/call  Ts/call  name
99.90      8.15      8.15         2      0.00    0.00    main
 0.00      8.15      0.00         2      0.00    0.00    diff

%           the percentage of the total running time of the
time        program used by this function.

cumulative  a running sum of the number of seconds accounted
seconds     for by this function and those listed above it.

self        the number of seconds accounted for by this
seconds     function alone.  This is the major sort for this
            listing.

calls       the number of times this function was invoked, if
            this function is profiled, else blank.

self        the average number of milliseconds spent in this
ms/call     function per call, if this function is profiled,
            else blank.

total       the average number of milliseconds spent in this
ms/call     function and its descendents per call, if this
            function is profiled, else blank.

name        the name of the function.  This is the minor sort
            for this listing.  The index shows the location of
            the function in the gprof listing.  If the index is
            in parenthesis it shows where it would appear in
            the gprof listing if it were to be printed.

Copyright (C) 2012 Free Software Foundation, Inc.

```

Figure 3: Profiling on HPC cluster using gprof

3.4 Optimization Strategy

We are not optimizing algorithm for matrix multiplication as here in this question we are supposed to do it in $\mathcal{O}(n^3)$ time complexity. We are optimizing little bits of code like writing $C[i][j] += A[i][k] * B[k][j]$ instead of $C[i][j] = C[i][j] + A[i][k] * B[k][j]$ to reduce the number of memory access.

The reason why $C[i][j] += A[i][k] * B[k][j]$ has better performance than $C[i][j] = C[i][j] + A[i][k] * B[k][j]$ is that it avoids an extra memory access for $C[i][j]$. In the first case, $C[i][j]$ is only read once and updated once. In the second case, $C[i][j]$ is read twice and updated once. This can make a difference when dealing with large matrices and cache memory.

3.5 Graph of Problem Size vs Algorithm Runtime

3.5.1 Graph of Problem Size vs Algorithm Runtime for LAB207 PCs




Figure 4: Mean Algorithm execution time vs Problem size plot for Lab 207 PC

3.5.2 Graph of Problem Size vs Runtime for HPC Cluster




Figure 5: Mean Algorithm execution time vs Problem size plot for Lab 207 PC

3.6 Graph of Problem Size vs End to End Runtime

3.6.1 Graph of Problem Size vs End to End Runtime for LAB207 PCs




Figure 6: Mean End to End execution time vs Problem size plot for Lab 207 PC

3.6.2 Graph of Problem Size vs End to End Runtime for HPC Cluster

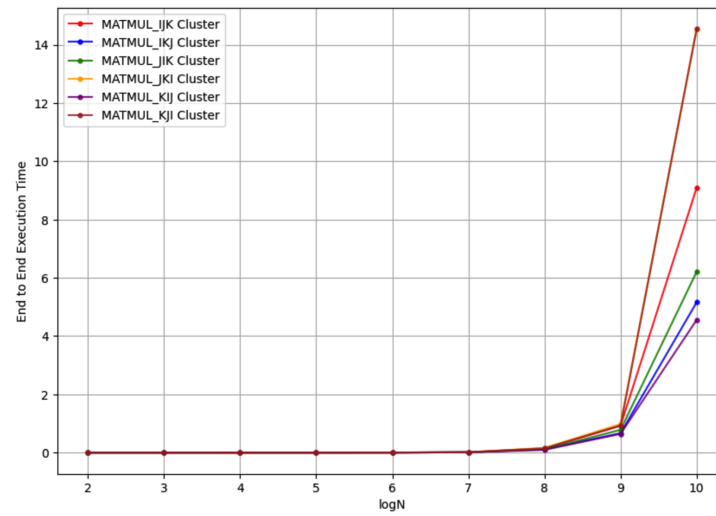


Figure 7: Mean End to End execution time vs Problem size plot for HPC Cluster

3.7 Discussion

The graph shows that runtime depends on the order of three loops i, j and k respectively.

The runtime was comparatively low for loop order KIJ and IKJ because of spatial locality of data. The code in innermost loop is simply $C[i][j] += A[i][k] * B[k][j]$; The order of loops in matrix multiplication affects spatial locality because it determines how efficiently the elements of the matrices are accessed in cache. The matrices are stored in row-major order, which means that consecutive elements in a row are stored next to each other in cache. This means that accessing elements in a row is faster than accessing elements in a column, because it reduces cache misses and increases cache hits.

Loop orders KIJ and IKJ have lower runtime because they access matrix A by rows and matrix B by columns. This minimizes cache misses and maximizes spatial locality. Loop orders KJI and JKI have higher runtime because they access both matrices by columns. This causes many cache misses and poor spatial locality.

- IJK: This loop order accesses matrix A by rows and matrix C by rows. It has good spatial locality for these matrices, but poor spatial locality for matrix B , which is accessed by columns. It has moderate performance.
- IKJ: This loop order accesses matrix A by rows and matrix B by columns. It has good spatial locality for both matrices, and also for matrix C , which is updated by rows. It has high performance.
- JIK: This loop order accesses matrix C by columns and matrix A by columns. It has poor spatial locality for these matrices, but good spatial locality for matrix B , which is accessed by rows. It has low performance.
- JKI: This loop order accesses matrix C by columns and matrix B by columns. It has poor spatial locality for both matrices, and also for matrix A , which is updated by rows. It has very low performance.
- KIJ: This loop order accesses matrix A by rows and matrix B by columns. It has good spatial locality for both matrices, and also for matrix C , which is updated by elements. It has high performance.
- KJI: This loop order accesses matrix B by columns and matrix C by columns. It has poor spatial locality for both matrices, but good spatial locality for matrix A , which is accessed by elements. It has very low performance.