
CS 301

High-Performance Computing

Lab 5 - B1

Matrix Multiplication using Transpose

Aditya Nawal (202001402)
Divya Patel (202001420)

March 23, 2023

Contents

1	Introduction	3
2	Hardware Details	3
2.1	Lab 207 PC	3
2.2	HPC Cluster	5
3	Problem B1	6
3.1	Description of the problem	6
3.2	Serial Complexity	6
3.3	Parallel Complexity	6
3.4	Profiling Information	7
3.5	Optimization Strategy	8
3.6	Graph of Problem Size vs Algorithm Runtime	8
3.6.1	Graph of Problem Size vs Algorithm Runtime for LAB207 PCs	8
3.6.2	Graph of Problem Size vs Algorithm Runtime for HPC Cluster	9
3.7	Graph of Problem Size vs End-to-End Runtime	9
3.7.1	Graph of Problem Size vs End to End Runtime for LAB207 PCs	9
3.7.2	Graph of Problem Size vs End to End Runtime for HPC Cluster	10
3.8	Discussion	10
4	Alternative Approach(Experiment)	11
4.1	Illustration of this method	11
4.1.1	Graph of Problem Size vs Algorithm Runtime	11

1 Introduction

This report investigates three different approaches to matrix multiplication. In Problem A-1, we explore the conventional matrix multiplication algorithm. In Problem B-1, we study the use of transpose matrix multiplication as a way to reduce the computational cost of matrix multiplication. Finally, in Problem C-1, we analyze the block matrix multiplication algorithm, which uses a divide and conquer strategy to compute the product of two large matrices.

Parallelization of the above algorithms can help reduce their computational cost and improve its performance. In this report, we present a detailed analysis of the parallelization process and evaluate their performance on matrices of different sizes. For each problem, we provide a detailed analysis of the algorithms, including their computational complexity.

2 Hardware Details

2.1 Lab 207 PC

- Architecture: x86_64
- CPU op-mode(s): 32-bit, 64-bit
- Byte Order: Little Endian
- CPU(s): 4
- On-line CPU(s) list: 0-3
- Thread(s) per core: 1
- Core(s) per socket: 4
- Socket(s): 1
- NUMA node(s): 1
- Vendor ID: GenuineIntel
- CPU family: 6
- Model: 60
- Model name: Intel(R) Core(TM) i5-4590 CPU @ 3.30GHz
- Stepping: 3
- CPU MHz: 3300.000
- CPU max MHz: 3700.0000
- CPU min MHz: 800.0000
- BogoMIPS: 6585.38

- Virtualization: VT-x
- L1d cache: 32K
- L1i cache: 32K
- L2 cache: 256K
- L3 cache: 6144K
- NUMA node0 CPU(s): 0-3
- Flags: fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm constant_tsc arch_perfmon pebs bts rep_good nopl xtopology nonstop_tsc aperfmperf eagerfpu pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 fma cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer aes xsave avx f16c rdrand lahf_lm abm epb invpcid_single tpr_shadow vnmi flexpriority ept vpid fsgsbase tsc_adjust bmi1 avx2 smep bmi2 erms invpcid xsaveopt dtherm ida arat pln pts

```
[student@localhost ~]$ ifconfig
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1 (Local Loopback)
    RX packets 60 bytes 5868 (5.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 60 bytes 5868 (5.7 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

p4p1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.100.64.86 netmask 255.255.255.0 broadcast 10.100.64.255
    inet6 fe80::b283:feff:fe97:d2f9 prefixlen 64 scopeid 0x20<link>
    ether b0:83:fe:97:d2:f9 txqueuelen 1000 (Ethernet)
    RX packets 32826 bytes 46075919 (43.9 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 8015 bytes 586362 (572.6 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

virbr0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 192.168.122.1 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:3a:16:71 txqueuelen 1000 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

Figure 1: IP address of Lab PC

2.2 HPC Cluster

- Architecture: x86_64
- CPU op-mode(s): 32-bit, 64-bit
- Byte Order: Little Endian
- CPU(s): 16
- On-line CPU(s) list: 0-15
- Thread(s) per core: 1
- Core(s) per socket: 8
- Socket(s): 2
- NUMA node(s): 2
- Vendor ID: GenuineIntel
- CPU family: 6
- Model: 63
- Model name: Intel(R) Xeon(R) CPU E5-2640 v3 @ 2.60GHz
- Stepping: 2
- CPU MHz: 1976.914
- BogoMIPS: 5205.04
- Virtualization: VT-x
- L1d cache: 32K
- L1i cache: 32K
- L2 cache: 256K
- L3 cache: 20480K
- NUMA node0 CPU(s): 0-7
- NUMA node1 CPU(s): 8-15

3 Problem B1

3.1 Description of the problem

Given two matrices A and B of dimensions $m \times n$ and $n \times p$, respectively, the standard matrix multiplication algorithm requires $O(mnp)$ operations to compute the product $C = AB$, where C is a matrix of dimensions $m \times p$. This computational cost is prohibitive for large matrices and, therefore, efficient algorithms for matrix multiplication are of great interest.

One possible approach to reduce the computational cost of matrix multiplication is to use the transpose of one of the matrices. In particular, we can compute the product $C = AB$ as follows:

$$C = A \cdot B^T, \tag{1}$$

where B^T is the transpose of B .

we convert $C_{ij} = \sum_{k=1}^n A_{ik}B_{kj}$ to $C_{ij} = \sum_{k=1}^n A_{ik}B_{jk}$ When you multiply two matrices A and B , you need to access the elements of B by columns, which are not stored contiguously in memory if you use row-major ordering. This can result in many cache misses and slow down the operation.

However, if you transpose B first, then you can access its elements by rows, which are stored contiguously in memory. This can reduce the number of cache misses and speed up the operation.

Of course, transposing B also requires accessing its elements non-sequentially, but this is done only once before the multiplication. The multiplication itself is done many times and benefits from cache locality.

3.2 Serial Complexity

This method requires same number of operations as the standard matrix multiplication algorithm, i.e.,

$$T_s = \mathcal{O}(n^3)$$

.

3.3 Parallel Complexity

$$T_p = \mathcal{O}\left(\frac{n^3}{p}\right)$$

where T_p is the parallel time complexity, N is the size of the matrices being multiplied and p is the number of processors.

This expression assumes that the workload is evenly distributed among all processors and that there are no overheads due to parallelization. In practice, however, there may be some overheads due to thread creation and synchronization which can affect the actual parallel time complexity.

3.4 Profiling Information

```

Flat profile:

Each sample counts as 0.01 seconds.
%   cumulative   self           self       total
time   seconds    seconds       calls   Ts/call   Ts/call   name
100.30      1.73      1.73           2        0.00      0.00   main
 0.00      1.73      0.00           2        0.00      0.00   diff

%           the percentage of the total running time of the
time        program used by this function.

cumulative  a running sum of the number of seconds accounted
seconds     for by this function and those listed above it.

self        the number of seconds accounted for by this
seconds     function alone. This is the major sort for this
           listing.

calls       the number of times this function was invoked, if
           this function is profiled, else blank.

self        the average number of milliseconds spent in this
ms/call     function per call, if this function is profiled,
           else blank.

total       the average number of milliseconds spent in this
ms/call     function and its descendents per call, if this
           function is profiled, else blank.

name        the name of the function. This is the minor sort
           for this listing. The index shows the location of
           the function in the gprof listing. If the index is
           in parenthesis it shows where it would appear in
           the gprof listing if it were to be printed.

Copyright (C) 2012-2014 Free Software Foundation, Inc.

Copying and distribution of this file, with or without modification,
are permitted in any medium without royalty provided the copyright
notice and this notice are preserved.

```

Figure 2: Screenshot of text file generated from profiling on Lab 207 PC using gprof

```

Flat profile:

Each sample counts as 0.01 seconds.
%   cumulative   self           self       total
time   seconds    seconds       calls   Ts/call   Ts/call   name
100.39      0.53      0.53           2        0.00      0.00   main
 0.00      0.53      0.00           2        0.00      0.00   diff

%           the percentage of the total running time of the
time        program used by this function.

cumulative  a running sum of the number of seconds accounted
seconds     for by this function and those listed above it.

self        the number of seconds accounted for by this
seconds     function alone. This is the major sort for this
           listing.

calls       the number of times this function was invoked, if
           this function is profiled, else blank.

self        the average number of milliseconds spent in this
ms/call     function per call, if this function is profiled,
           else blank.

total       the average number of milliseconds spent in this
ms/call     function and its descendents per call, if this
           function is profiled, else blank.

name        the name of the function. This is the minor sort
           for this listing. The index shows the location of
           the function in the gprof listing. If the index is
           in parenthesis it shows where it would appear in
           the gprof listing if it were to be printed.

Copyright (C) 2012-2014 Free Software Foundation, Inc.

```

Figure 3: Profiling on HPC cluster using gprof

3.5 Optimization Strategy

The given code implements an optimization strategy for matrix multiplication based on parallelization using OpenMP. The outer two loops are parallelized using the `collapse(2)` clause and the innermost loop uses a reduction to calculate the sum.

The `collapse(2)` clause combines the two nested loops into a single loop that is then parallelized across threads. This allows for more efficient workload distribution among threads and can potentially improve performance.

The `reduction(+:sum)` clause is used for the innermost loop to sum the product of matrices A and B_transpose and store the result in the variable sum. The reduction clause ensures that the final result is the sum of the individual results obtained by each thread, thus avoiding race conditions and ensuring the correctness of the calculation.

In summary, this optimization strategy improves performance by parallelizing the computation using OpenMP. The outer two loops are collapsed into a single loop that is parallelized across threads and a reduction is used in the innermost loop to calculate the sum. This can potentially improve performance by distributing workload among multiple threads and reducing overall execution time.

3.6 Graph of Problem Size vs Algorithm Runtime

3.6.1 Graph of Problem Size vs Algorithm Runtime for LAB207 PCs

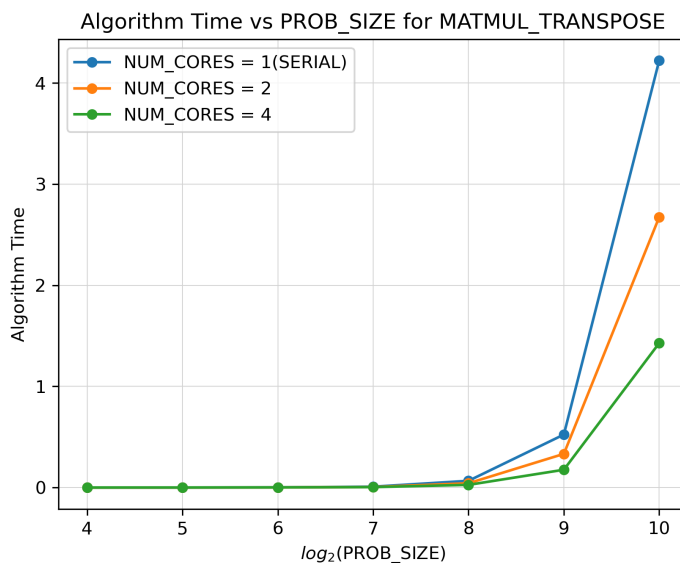


Figure 4: Graph of Problem Size vs Algorithm Runtime for Lab PC

3.6.2 Graph of Problem Size vs Algorithm Runtime for HPC Cluster




Figure 5: Graph of Problem Size vs Algorithm Runtime for HPC cluster

3.7 Graph of Problem Size vs End-to-End Runtime

3.7.1 Graph of Problem Size vs End to End Runtime for LAB207 PCs




Figure 6: Graph of Problem Size vs End-to-End Runtime for Lab PC

3.7.2 Graph of Problem Size vs End to End Runtime for HPC Cluster




Figure 7: Graph of Problem Size vs End-to-End Runtime for HPC cluster

3.8 Discussion

In this report, we have explored the idea of transposing a matrix before multiplication as a way to improve the computational efficiency of the operation. We have seen that transposing a matrix can improve the cache locality of the data access and reduce the number of cache misses. We have also learned that transposing a matrix has some properties such as:

1. The transpose of a scalar multiple of a matrix is equal to the scalar multiple of the transpose of the matrix.
2. The transpose of the product of two matrices is equal to the product of their transposes in reverse order.

These properties can help us manipulate matrices and simplify expressions involving transposes. However, we should note that transposing a matrix before multiplication is not always faster than normal matrix multiplication, as it depends on various factors such as the size and shape of the matrices, the implementation of the algorithm, and the hardware architecture.

Also, parallelizing the code using OpenMP can improve the performance of the algorithm by distributing the workload among multiple threads. However, the improvement in performance decreases as the number of processors increases. This is because the overheads due to parallelization increase as the number of processors increases.

4 Alternative Approach(Experiment)

Consider two matrices A and B with dimensions $m \times n$ and $n \times p$, respectively. The product of these matrices is a matrix C with dimensions $m \times p$. To compute the element c_{ij} of matrix C , we need to compute the dot product of the i^{th} row of matrix A and the j^{th} column of matrix B .

If we store matrix A in a row-major 1-D array and matrix B in a column-major 1-D array, then accessing the elements required for computing the dot product becomes more efficient. This is because consecutive elements in a row or column are stored in consecutive memory locations, improving cache locality.

4.1 Illustration of this method

Consider two matrices A and B with dimensions 2×3 and 3×2 , respectively:

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \quad B = \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix}$$

We can convert matrix A to a row-major 1-D array and matrix B to a column-major 1-D array as follows:

$$A_{row_major} = [1, 2, 3, 4, 5, 6] \quad B_{col_major} = [7, 9, 11, 8, 10, 12]$$

Now let's compute the element $c_{1,1}$ of the resulting matrix C . This element is obtained by computing the dot product of the first row of matrix A and the first column of matrix B . In our row-major and column-major representations this corresponds to:

$$\begin{aligned} c_{1,1} &= A_{row_major}[0] \cdot B_{col_major}[0] + A_{row_major}[1] \cdot B_{col_major}[1] + A_{row_major}[2] \cdot B_{col_major}[2] \\ &= (1 \cdot 7) + (2 \cdot 9) + (3 \cdot 11) \\ &= 58 \end{aligned}$$

4.1.1 Graph of Problem Size vs Algorithm Runtime

From the graph we can see that the performance of this method is even better than the block matrix multiplication method for this problem size range.

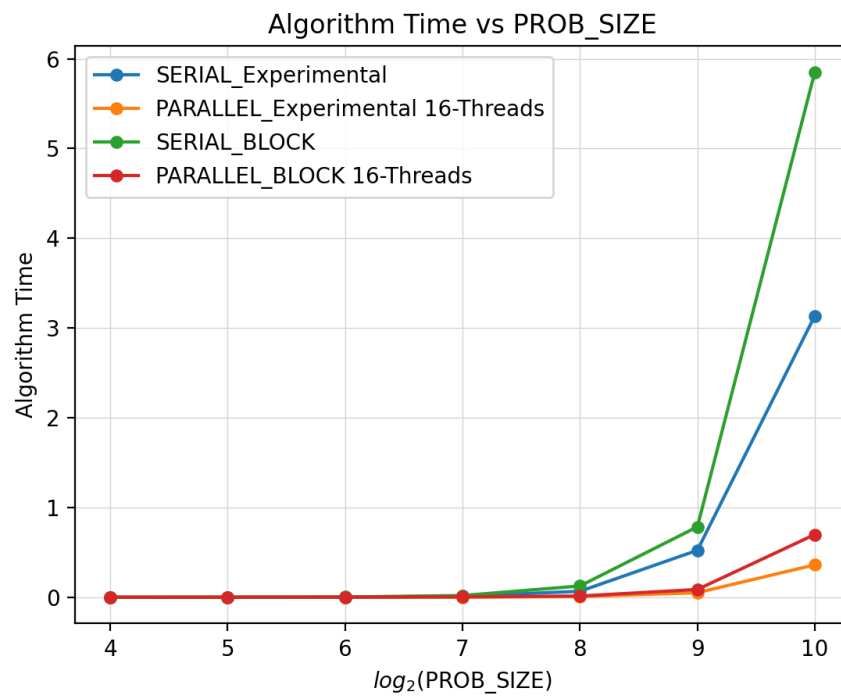


Figure 8: Graph of Problem Size vs Algorithm Runtime for HPC Cluster